

Kernels I

ENCE361 Embedded Systems 1

Course Coordinator: Ciaran Moore (ciaran.moore@Canterbury.ac.nz)

Lecturer: Le Yang (le.yang@canterbury.ac.nz)

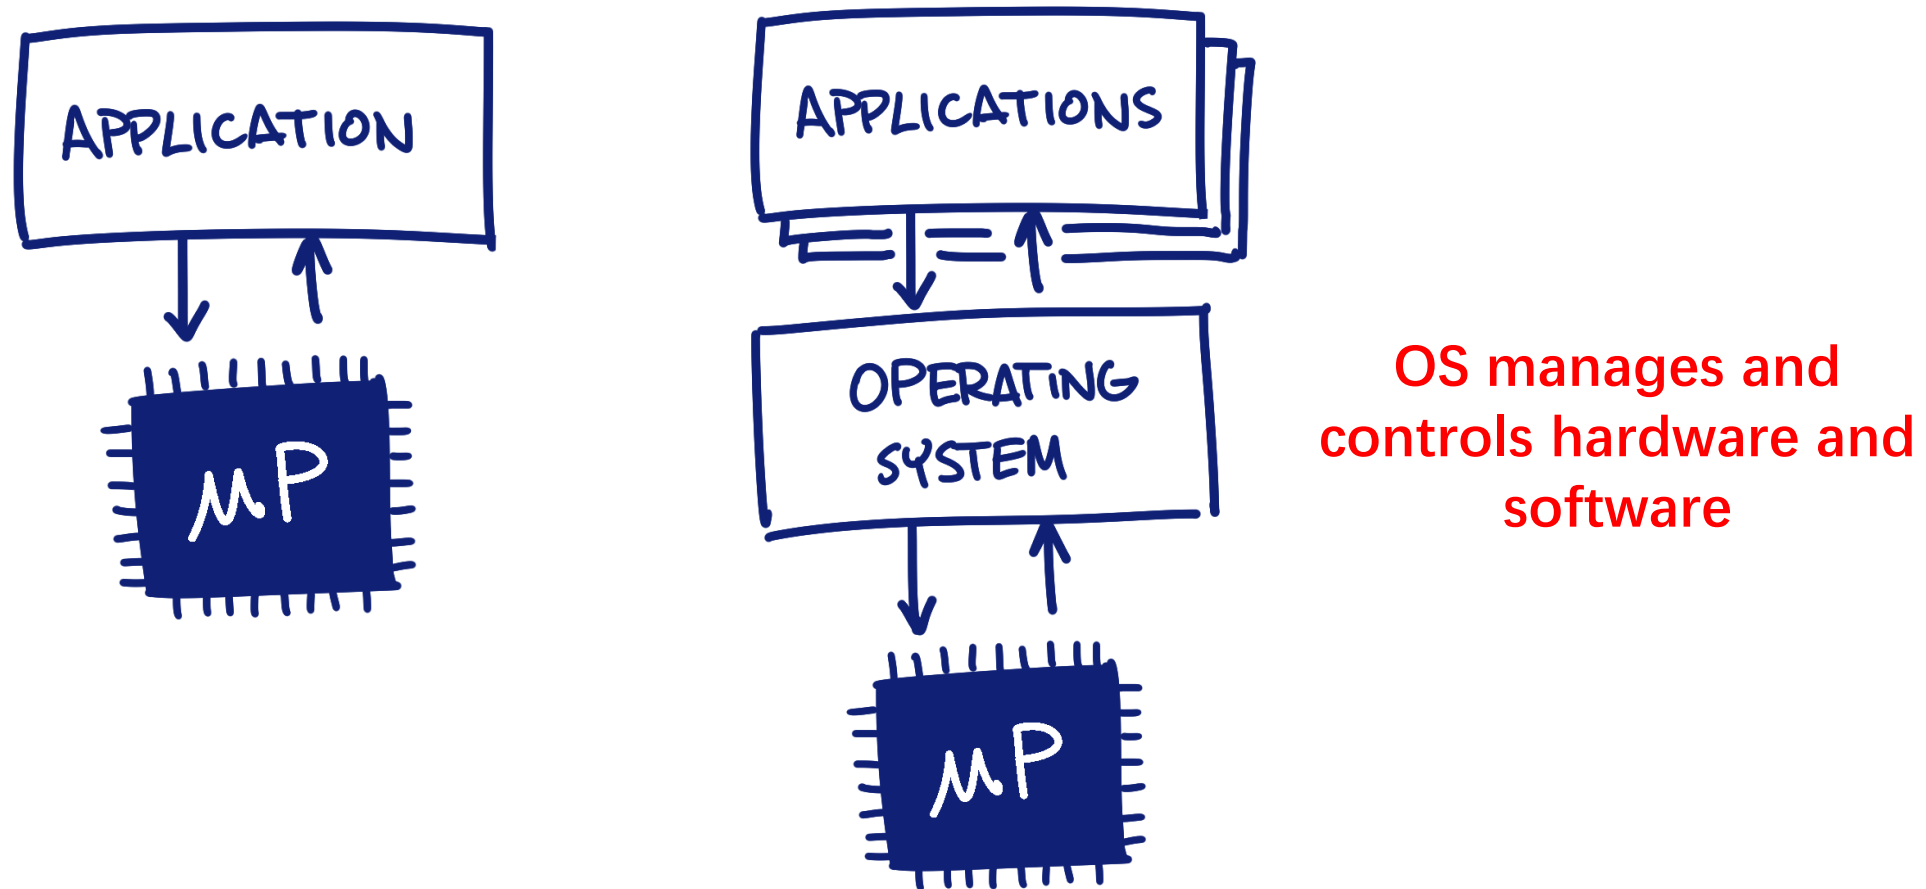
Department of Electrical and Computer Engineering

Where we're going today

- **Basics of kernels**
- Round-robin, time-triggered and interrupt-triggered schedulers
- Design example
- Homework

Kernel Basics (1)



- Embedded software vs. most other software



Kernel Basics (2)

- Complex embedded software

```
void main (void)
{
    while (true)
    {
        // Do everything!
        
    }
}
```



- What is a **kernel**?
 - Main component of an OS
 - Bridge between applications and actual processing done at hardware
 - Responsibility: managing system resource

Functions of Kernels

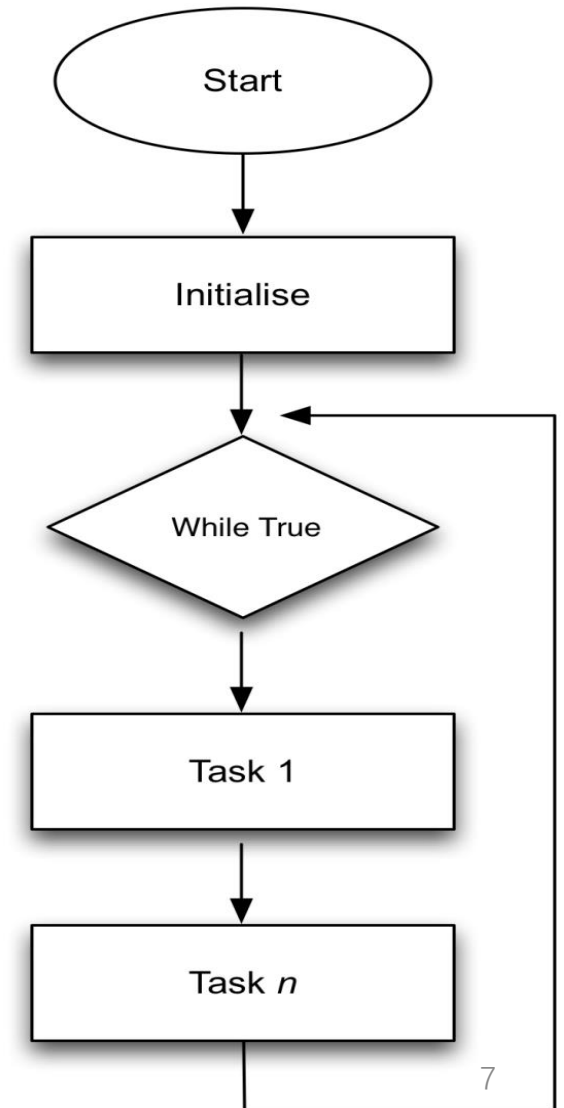
- What does a kernel do?
 - Schedule background tasks → **scheduler**
 - All processing is achieved in a timely manner
 - All tasks are completed reliably
 - May ensure that higher priority tasks are completed first
 - May allow higher priority tasks to preempt lower priority tasks
 - May provide the ability for tasks to communicate (inter-task communication)
 - May provide the ability for resources (e.g., memory blocks, peripherals, etc.) to be shared

Where we're going today

- Basics of kernels
- **Round-robin, time-triggered and interrupt-triggered schedulers**
- Design example
- Homework

Round-Robin Scheduler (1)

- Round-robin scheduler (c.f., [polling](#))
 - Free-running cyclic
- **Advantages** 😊
 - Simple to implement
 - Same worst-case response time for every task
- **Disadvantages** ☹️
 - Same worst-case response time for every task
 - All tasks have the same priority level
 - Tasks cannot be very long
 - Time slicing may be used to split a task into smaller chunks



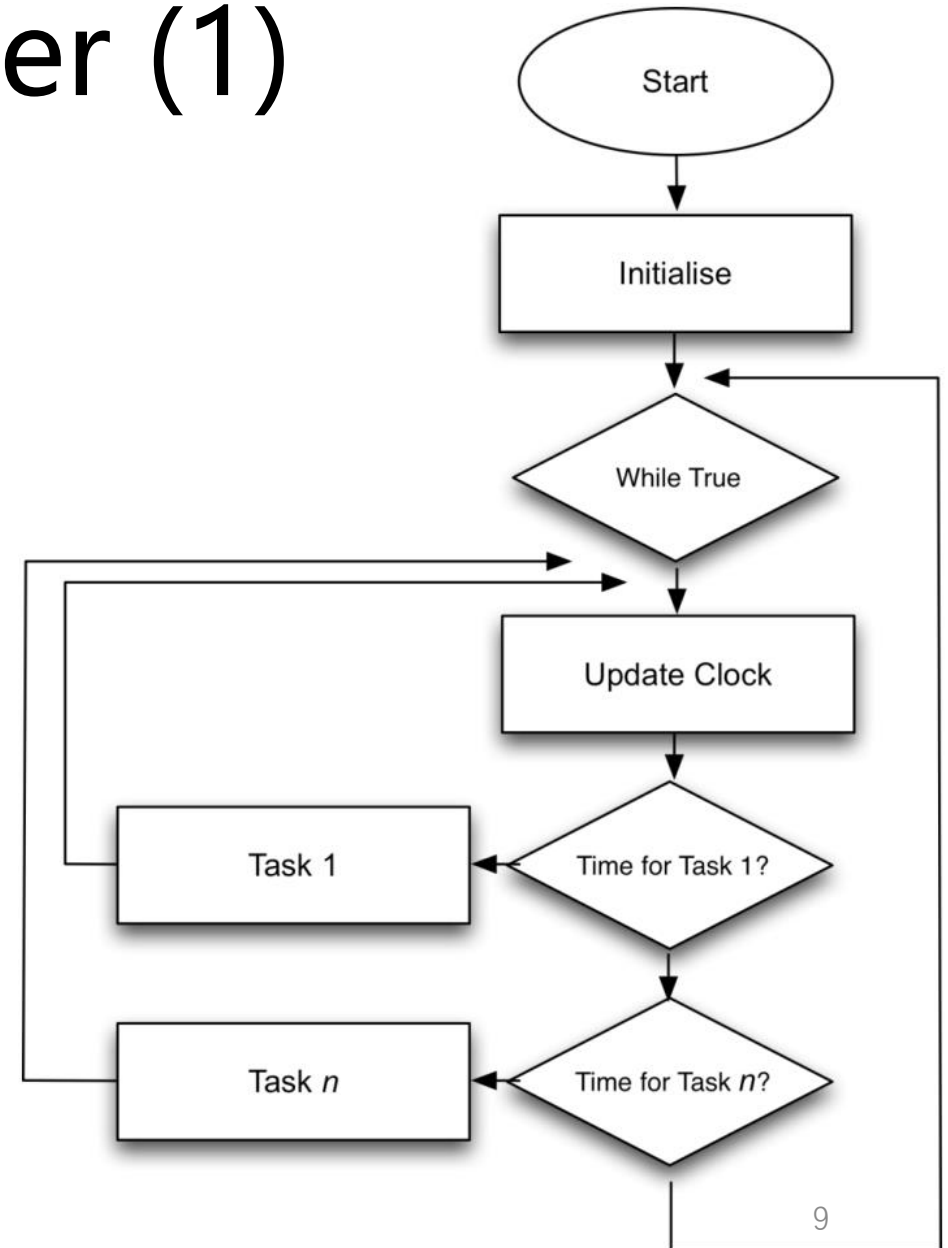
Round-Robin Scheduler (2)

```
void main (void)
{
    while (true)
    {
        if (/* I/O device A needs service */)
        {
            // Handle data to or from Device A
        }
        if (/* I/O device B needs service */)
        {
            // Handle data to or from Device B
        }
        . . . . .

        if (/* I/O device Z needs service */)
        {
            // Handle data to or from Device Z
        }
    }
}
```


Time-Triggered Scheduler (1)

- Time-triggered scheduler
 - 'Time-division multiplexing' of MCU
- Advantages 😊
 - Perform tasks at controlled times
 - Limited task prioritization
- Disadvantages ☹️
 - Worst-case response time depends on scheduled task run times
 - Tasks cannot be very long
 - Time slicing may be used to split a task into smaller chunks



Time-Triggered Scheduler (2)

```
void timerInterruptHandler(void) {  
    for (uint8_t i = 0; i < N_TASKS; i++) {  
        if (scheduled_task[i].delay == 0) {  
            scheduled_task[i].ready = true;  
            scheduled_task[i].delay = scheduled_task[i].period; // Reset task  
        } else {  
            scheduled_task[i].delay -= 1;  
        }  
    }  
}
```

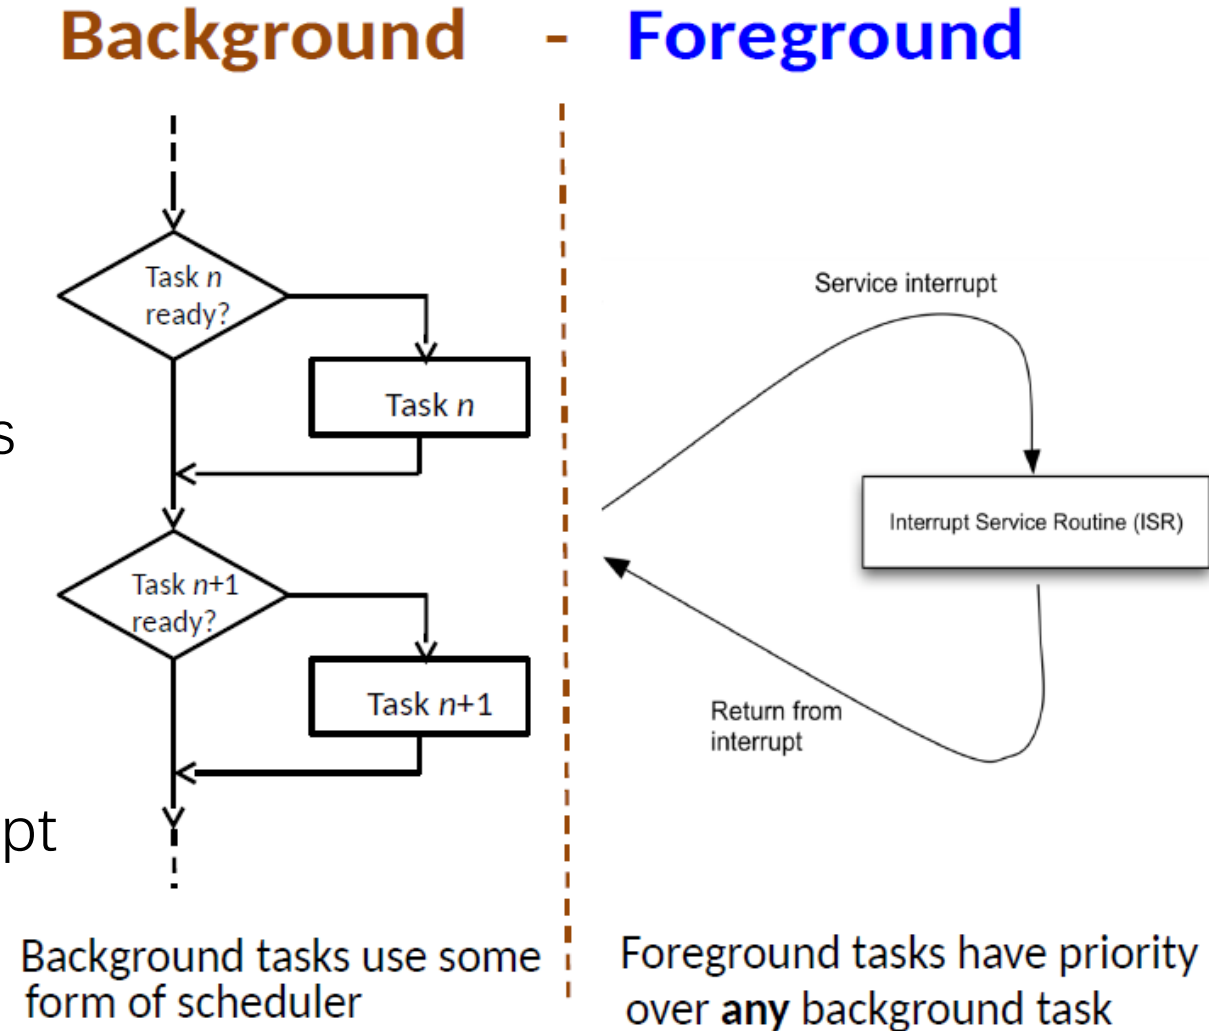
// Please, do not code like this ...

```
void main (void) {  
    while (true) {  
        for (uint8_t i = 0; i < N_TASKS; i++) {  
            if (scheduled_task[i].ready) {  
                // Call task handler  
                ...  
                scheduled_task[i].ready = false;  
            }  
        }  
    }  
}
```

// Please, do not code like this ...

Interrupt-Triggered Scheduler (1)

- Interrupt-triggered scheduler
 - Foreground/background modeling
- Advantages 😊
 - Fast response to asynchronous events
 - Interrupt prioritization
- Disadvantages ☹️
 - Worst-case response time of background tasks depends on interrupt arrival rates
 - Background tasks cannot take 'very long'



Interrupt-Triggered Scheduler (2)


```
volatile boolean flagDeviceA = false; . . . ; flagDeviceZ = false;
void deviceAIntHandler (void) { /* Handle Device A I/O, set flag */ }
. . . . .
void deviceZIntHandler (void) { /* Handle Device Z I/O, set flag */ }
void main (void)
{
    . . . . .
    while (true)
    {
        if (flagDeviceA)
        {
            // Process data to or from Device A and reset flag
        }
        . . . . .
        if (flagDeviceZ)
        {
            // Process data to or from Device Z and reset flag
        }
    }
}
```

Where we're going today

- Basics of kernels
- Round-robin, time-triggered and interrupt-triggered schedulers
- **Design example**
- Homework

Anti-Lock Braking (ABS) System

- Accurately monitor speeds of four wheels to detect potential skid
- Interrupts for real-time clock and wheel pulses
- Pulse rate up to 1320 pulses/s (200 km/h)
 - $1320 \times 4 = 5280$ pulses/s
- Wheel speed evaluated every 5 pulses
 - 264 speed estimates/s for each wheel



Tasks

Task	Function	Required frequency	Execution time
Real-time clock (RTC)	ISR_SysTick ()	$10^5 / \text{s}$	$2.5 \mu\text{s}$
Wheel pulses	ISR_Wheels ()	$5280 / \text{s}$	$20 \mu\text{s}$
Task A: FLwheel speed	speedFL ()	$264 / \text{s}$	$30 \mu\text{s}$
Task B: FRwheel speed	speedFR ()	$264 / \text{s}$	$30 \mu\text{s}$
Task C: RLwheel speed	speedRL ()	$264 / \text{s}$	$30 \mu\text{s}$
Task D: RRwheel speed	speedRR ()	$264 / \text{s}$	$30 \mu\text{s}$
Task E: Monitor relative speeds	monitorWheels ()	$200 / \text{s}$	$250 \mu\text{s}$
Task F: Monitor driver actions	monitorDriver ()	$50 / \text{s}$	1.5 ms
Task G: Activate ABS	activateABS ()	Infrequent	5 ms
Task H: Display status	displayABS ()	$20 / \text{s}$	7.5 ms

Round-Robin Kernel-based Design

```
void main (void)
{
    // Initialise
    ....

    // Start round-robin kernel
    while (1)
    {
        speedFL ();           // 30 µs, 264/s. True frequency < 106/s (insufficient!)
        speedFR ();
        speedRL ();
        speedRR ();
        if (monitorWheels ()) // 250 µs, 200/s. True frequency < 106/s (insufficient!)
            activateABS ();
        monitorDriver ();     // 1.5 ms, 50/s. True frequency < 106/s (not necessary)
        displayABS ();        // 7.5ms, 20/s. True frequency < 106/s (not necessary)
    }
}
```


Time-Triggered Kernel-based Design

```
void main (void)
{
// Initialise
....
// Start time-triggered kernel
while (1)
{
    if (flagFL) { speedFL (); flagFL = 0; }           // 30 µs, 264/s = True frequency
    if (flagFR) { speedFR (); flagFR = 0; }
    if (flagRL) { speedRL (); flagRL = 0; }
    if (flagRR) { speedRR (); flagRR = 0; }
    if (monitorWheels ())                             // 250 µs, 200/s = True frequency
        activateABS ();
    if (monitorDriverFlag) { monitorDriver (); monitorDriverFlag = 0;} // 1.5 ms, 50/s = True frequency
    if (displayABSFlag) { displayABS (); displayABSFlag = 0;}         // 7.5ms, 20/s = True frequency
}
}
```

Homework

1. What is the major difference between the time-triggered and interrupt-triggered schedulers?
2. What determines the resolution required for the real-time clock in the implementation of the ABS braking system?
3. In the time-triggered version of the ABS system, what is the maximum amount of time that could be spent on executing a chunk of the monitorDriver and displayABS tasks without causing the deadlines for wheel speed processing to be missed?